KPIS Project: Timeline

Topic: Document Classification Performance Comparison Study (Encoder Architecture)
Focus: How LoRA fine-tuning settings affect a BERT encoder's document classification performance
If there is time: out of interest, an exploratory look inside the fine-tuned models (vector space, explainable AI)
Authors: Muhammad and Sabrina

Weekly plan (winter semester 2026/27)

Week from	Activity	Deadline / event
27.10	Research review, plan settings	27.10 milestones
03.11	Research review, plan settings, paper outline	03.11 paper outline
10.11	Setup: notebooks on Kaggle, data split, first baseline run (approach 1)	
17.11	Run experiments: LoRA study, first half of the settings (approach 2)	
24.11	Run experiments: LoRA study, second half of the settings (approach 2)	
01.12	Run experiments: second dataset (approach 3), slides for the interim talk	~02./03.12 interim talk
08.12	Evaluation: tables and plots (+ approach 5, if time)	
15.12	Evaluation: what worked, what didn't (+ approach 5, if time)	
21.12	Write up (Christmas break)	
05.01	Write up, review others' papers	05.01 draft paper
12.01	Revise paper, prepare talk	12.01 reviews
19.01	Prepare talk	19.01 final paper
20.01	Talk	20./21.01 or 27./28.01

Approaches

1. **Baseline:** BERT as it is, simple classifier on top
2. **LoRA study:** change one setting at a time, look at accuracy and cost, every setting 3 times
3. **Second dataset:** short sentences vs. longer documents, do the same settings still work?
4. **Fallback:** just inference, compare pretrained encoders without any LoRA/fine-tuning
5. **If time:** exploratory look inside the trained models to explain the results (vector space, explainable AI)

Dates are planned and may shift slightly as the work progresses.
Table and PDF layout generated with Claude Code. Structure and content about 95% planned by us, with some suggestions from Claude Code (e.g. 3 runs per setting) and help with refactoring grammar and phrasing.